

Preparing FastAPI for Production: A Comprehensive Guide

25 min read · Oct 7, 2024



Raman Bazhanau

Follow



Listen



Share



Table of Contents

1. Introduction to FastAPI and Production Deployment

- What is FastAPI?

- **The Importance of Proper Production Deployment**
- **What We'll Cover in This Guide**

2. Configuring FastAPI for Production

- **Development vs. Production Configurations**
- **Using Environment Variables**
- **Configuration Management Best Practices**
- **A Production-Ready FastAPI Configuration Example**

3. Running FastAPI in Production

- **ASGI Servers for FastAPI (Uvicorn, Hypercorn, Gunicorn)**
- **Process Management and Daemonization**
- **Best Practices for Running FastAPI in Production**

4. Security Considerations

- **HTTPS and SSL/TLS Configuration**
- **API Authentication and Authorization**
- **CORS (Cross-Origin Resource Sharing) Setup**
- **Rate Limiting Implementation**

5. Performance Optimization

- **Async and Await Usage in FastAPI**
- **Database Connection Pooling**
- **Caching Strategies (In-memory, Redis)**
- **Request Validation and Response Serialization Optimization**

6. Logging and Monitoring

- **Setting Up Proper Logging in FastAPI**
- **Log Rotation and Management**
- **Monitoring Tools and Integration**
- **Useful Metrics to Track**

7. Containerization with Docker

- Benefits of Using Docker with FastAPI
- Sample Dockerfile for a FastAPI Application
- Docker Compose for Multi-Container Setups
- Best Practices for Docker in Production

8. Deployment Strategies

- Deployment Options (VPS, PaaS, Serverless)
- CI/CD Pipelines for FastAPI Applications
- Blue-Green Deployments and Rolling Updates
- Deployment Scripts

9. Scaling FastAPI Applications

- Horizontal vs. Vertical Scaling
- Load Balancing Options
- Database Scaling Considerations
- Serverless Scaling with FastAPI

10. Conclusion and Best Practices

- Key Takeaways
- Best Practices Checklist
- Final Thoughts

1. Introduction to FastAPI and Production Deployment

FastAPI has quickly become one of the most popular Python web frameworks for building APIs. Its combination of speed, simplicity, and powerful features makes it an excellent choice for developers looking to create robust and efficient web applications. However, moving from development to production requires careful consideration and proper configuration to ensure your FastAPI application performs optimally and securely in a real-world environment.

In this comprehensive guide, we'll explore the essential steps and best practices for preparing your FastAPI application for production deployment. Whether you're a

seasoned developer or just getting started with FastAPI, this article will provide you with valuable insights and practical advice to help you confidently deploy your application.

What is FastAPI?

FastAPI is a modern, fast (high-performance) web framework for building APIs with Python 3.6+ based on standard Python type hints. It's designed to be easy to use, fast to code, ready for production, and suitable for building robust and scalable applications. Key features of FastAPI include:

1. *Fast execution*: FastAPI is one of the fastest Python frameworks available, thanks to its use of Starlette and Pydantic.
2. *Easy to use*: With intuitive syntax and excellent documentation, FastAPI has a gentle learning curve.
3. *Built-in validation*: FastAPI uses Pydantic for data validation, serialization, and documentation.
4. *Automatic API documentation*: OpenAPI (Swagger) and ReDoc documentation are generated automatically.
5. *Python type hints*: FastAPI leverages Python's type hinting system for better code quality and IDE support.

The Importance of Proper Production Deployment

While FastAPI makes it easy to create APIs quickly, deploying an application to production requires additional considerations. A production environment faces real-world challenges such as high traffic, potential security threats, and the need for scalability and reliability. Proper production deployment ensures that your FastAPI application:

1. Performs efficiently under load
2. Remains secure against potential threats
3. Scales to meet growing demand
4. Provides a stable and reliable service to users
5. Can be easily monitored and maintained

What We'll Cover in This Guide

In this comprehensive guide, we'll walk through the essential aspects of preparing your FastAPI application for production deployment. We'll cover:

1. Configuring FastAPI for production environments
2. Running FastAPI with production-grade ASGI servers
3. Implementing crucial security measures
4. Optimizing performance for high-traffic scenarios
5. Setting up logging and monitoring
6. Containerizing your application with Docker
7. Exploring deployment strategies and scaling options

By the end of this guide, you'll have a solid understanding of what it takes to deploy a FastAPI application in a production environment, along with practical examples and best practices to follow.

2. Configuring FastAPI for Production

When moving your FastAPI application from development to production, proper configuration is crucial. This section will explore the key differences between development and production configurations, and provide practical advice on setting up your FastAPI app for a production environment.

Development vs. Production Configurations

In development, you typically prioritize convenience and debugging capabilities. However, in production, the focus shifts to security, performance, and reliability. Here are some key differences:

1. Debug mode: Disabled in production to prevent exposing sensitive information.
2. Error handling: Detailed errors in development, generic error messages in production.
3. Performance optimizations: Often bypassed in development, crucial in production.
4. Security features: Might be relaxed in development, must be strictly enforced in production.

Using Environment Variables

Environment variables are a crucial tool for managing configuration in production environments. They allow you to:

1. Keep sensitive information out of your codebase
2. Easily change configuration without modifying code
3. Use different settings for different environments (dev, staging, production)

Here's an example of how to use environment variables in your FastAPI application:

```
import os
from fastapi import FastAPI
from pydantic_settings import BaseSettings

class Settings(BaseSettings):
    app_name: str = "MyFastAPI App"
    admin_email: str
    database_url: str
    secret_key: str

    class Config:
        env_file = ".env"

settings = Settings()
app = FastAPI()

@app.get("/info")
async def info():
    return {
        "app_name": settings.app_name,
        "admin_email": settings.admin_email,
        "database_url": settings.database_url[:10] + "...", # Truncate for security
    }
```

In this example, we use `pydantic_settings` to create a `Settings` class that loads configuration from environment variables. The `env_file = ".env"` line allows loading variables from a `.env` file, which is useful for development but should be avoided in production (use actual environment variables instead).

Configuration Management Best Practices

1. *Use a configuration management tool:* Tools like Ansible, Puppet, or Chef can help manage configurations across multiple environments.
2. *Implement a secrets management system:* Use tools like HashiCorp Vault or AWS Secrets Manager to securely store and manage sensitive information.
3. *Use different configurations for different environments:* Maintain separate configuration files or environment variable sets for development, staging, and production.
4. *Version your configurations:* Keep your configuration files in version control, but ensure that sensitive data is not included.
5. *Use a centralized configuration service:* For complex, distributed systems, consider using tools like etcd or Consul for centralized configuration management.

A Production-Ready FastAPI Configuration Example

Here's a more comprehensive example of a production-ready FastAPI configuration:

```
import os
from fastapi import FastAPI, HTTPException
from pydantic_settings import BaseSettings
from functools import lru_cache

class Settings(BaseSettings):
    app_name: str = "MyFastAPI App"
    admin_email: str
    database_url: str
    secret_key: str
    allowed_hosts: list = ["*"]
    debug: bool = False

    class Config:
        env_file = ".env"

@lru_cache()
def get_settings():
    return Settings()

app = FastAPI()

@app.middleware("http")
async def validate_host(request, call_next):
```

```
settings = get_settings()
host = request.headers.get("host", "").split(":")[0]
if settings.debug or host in settings.allowed_hosts:
    return await call_next(request)
raise HTTPException(status_code=400, detail="Invalid host")


@app.get("/info")
async def info():
    settings = get_settings()
    return {
        "app_name": settings.app_name,
        "admin_email": settings.admin_email,
        "debug_mode": settings.debug
    }

if __name__ == "__main__":
    import uvicorn
    uvicorn.run(app, host="0.0.0.0", port=8000)
```

This configuration includes:

1. Environment-based settings with `pydantic_settings`
2. Host validation middleware for security
3. Caching of settings with `lru_cache` for performance
4. A simple info endpoint that respects the debug setting

To visualize the flow of configuration in a FastAPI application, here's a diagram:



This setup provides a solid foundation for a production-ready FastAPI application, balancing security, performance, and flexibility.

3. Running FastAPI in Production

When it comes to running FastAPI in a production environment, choosing the right ASGI (Asynchronous Server Gateway Interface) server is crucial. In this section, we'll explore different ASGI servers, their pros and cons, and how to run FastAPI with each of them.

ASGI Servers for FastAPI

FastAPI, being an ASGI framework, requires an ASGI server to run. The most common options are:

1. Uvicorn
2. Hypercorn
3. Gunicorn (with Uvicorn workers)

Let's examine each of these options in detail.

1. Uvicorn

Uvicorn is a lightning-fast ASGI server implementation, using uvloop and httptools for optimal performance.

Pros:

- Very fast and lightweight
- Easy to use and configure
- Supports HTTP/1.1 and WebSockets

Cons:

- Limited built-in process management
- Doesn't support HTTP/2 out of the box

To run FastAPI with Uvicorn:

```
uvicorn main:app --host 0.0.0.0 --port 8000 --workers 4
```

In this command:

- `main:app` refers to the `app` object in your `main.py` file
- `--host 0.0.0.0` allows external access
- `--port 8000` sets the port number
- `--workers 4` creates 4 worker processes

2. Hypercorn

Hypercorn is another ASGI server that supports HTTP/1, HTTP/2, and WebSockets.

Pros:

- Supports HTTP/2 and WebSockets
- Good performance

- More configuration options than Uvicorn

Cons:

- Slightly more complex to set up
- May be slower than Uvicorn for HTTP/1.1

To run FastAPI with Hypercorn:

```
hypercorn main:app --bind 0.0.0.0:8000 --workers 4
```

3. Gunicorn with Uvicorn workers

Gunicorn is a robust, production-ready server that can use Uvicorn workers to run FastAPI applications.

Pros:

- Production-ready with advanced features
- Excellent process management
- Can leverage Uvicorn's speed

Cons:

- More complex setup
- Requires additional configuration

To run FastAPI with Gunicorn and Uvicorn workers:

```
gunicorn main:app --workers 4 --worker-class uvicorn.workers.UvicornWorker --bi
```

Process Management and Daemonization

For long-running production deployments, you'll want to ensure your FastAPI application runs continuously and starts automatically if the server reboots.

1. Using Systemd

On many Linux systems, you can use systemd to manage your FastAPI application as a service. Here's an example systemd service file:

```
[Unit]
Description=FastAPI application
After=network.target

[Service]
User=youruser
WorkingDirectory=/path/to/your/app
ExecStart=/path/to/your/venv/bin/gunicorn main:app --workers 4 --worker-class uvicorn
Restart=always

[Install]
WantedBy=multi-user.target
```

Save this file as `/etc/systemd/system/fastapi.service`, then enable and start the service:

```
sudo systemctl enable fastapi
sudo systemctl start fastapi
```

2. Using Supervisor

Supervisor is another popular tool for process management. Here's an example configuration:


```
[program:fastapi]
command=/path/to/your/venv/bin/gunicorn main:app --workers 4 --worker-class uvicorn
directory=/path/to/your/app
user=youruser
autostart=true
autorestart=true
stderr_logfile=/var/log/fastapi.err.log
stdout_logfile=/var/log/fastapi.out.log
```

Save this as `/etc/supervisor/conf.d/fastapi.conf` , then update and start the service:

```
sudo supervisorctl reread
sudo supervisorctl update
sudo supervisorctl start fastapi
```

Best Practices for Running FastAPI in Production

1. **Use multiple workers:** The number of workers should generally be 2–4 times the number of CPU cores.
2. **Implement proper logging:** Configure your ASGI server to log to appropriate files.
3. **Use a reverse proxy:** Place Nginx or Apache in front of your FastAPI application for additional features and security.
4. **Monitor your application:** Use tools like Prometheus and Grafana to keep track of your application's health and performance.
5. **Implement graceful shutdowns:** Configure your server to handle shutdowns gracefully to prevent request interruptions.



This diagram will help visualize how all the components we've discussed fit together in a production environment.

4. Security Considerations

When deploying a FastAPI application to production, security should be a top priority. This section will cover essential security measures to protect your application and its users.

HTTPS and SSL/TLS Configuration

Implementing HTTPS is crucial for encrypting data in transit and ensuring the integrity of your API.

Setting up HTTPS

1. Obtain an SSL/TLS certificate:

- Use a service like Let's Encrypt for free certificates
- Purchase a certificate from a trusted Certificate Authority for commercial applications

2. Configure your reverse proxy (e.g., Nginx) to handle SSL/TLS:

```
server {  
    listen 443 ssl;  
    server_name yourdomain.com;  
  
    ssl_certificate /path/to/your/certificate.crt;  
    ssl_certificate_key /path/to/your/certificate.key;  
  
    ssl_protocols TLSv1.2 TLSv1.3;  
    ssl_prefer_server_ciphers on;  
    ssl_ciphers ECDHE-ECDSA-AES128-GCM-SHA256:ECDHE-RSA-AES128-GCM-SHA256:ECDHE
```

```
    location / {  
        proxy_pass http://localhost:8000;  
        proxy_set_header Host $host;  
        proxy_set_header X-Real-IP $remote_addr;  
    }  
}
```

3. Redirect HTTP to HTTPS:

```
server {  
    listen 80;  
    server_name yourdomain.com;  
    return 301 https://$server_name$request_uri;  
}
```

API Authentication and Authorization

Implementing proper authentication and authorization is essential for protecting your API endpoints.

1. JWT Authentication

JSON Web Tokens (JWT) are a popular choice for API authentication. Here's a basic example of implementing JWT authentication in FastAPI:

```
from fastapi import FastAPI, Depends, HTTPException, status  
from fastapi.security import OAuth2PasswordBearer, OAuth2PasswordRequestForm  
from jose import JWTError, jwt
```

```

from passlib.context import CryptContext
from datetime import datetime, timedelta

SECRET_KEY = "your-secret-key"
ALGORITHM = "HS256"
ACCESS_TOKEN_EXPIRE_MINUTES = 30

app = FastAPI()
pwd_context = CryptContext(schemes=["bcrypt"], deprecated="auto")
oauth2_scheme = OAuth2PasswordBearer(tokenUrl="token")

def create_access_token(data: dict):
    to_encode = data.copy()
    expire = datetime.utcnow() + timedelta(minutes=ACCESS_TOKEN_EXPIRE_MINUTES)
    to_encode.update({"exp": expire})
    encoded_jwt = jwt.encode(to_encode, SECRET_KEY, algorithm=ALGORITHM)
    return encoded_jwt

async def get_current_user(token: str = Depends(oauth2_scheme)):
    credentials_exception = HTTPException(
        status_code=status.HTTP_401_UNAUTHORIZED,
        detail="Could not validate credentials",
        headers={"WWW-Authenticate": "Bearer"},
    )
    try:
        payload = jwt.decode(token, SECRET_KEY, algorithms=[ALGORITHM])
        username: str = payload.get("sub")
        if username is None:
            raise credentials_exception
    except JWTError:
        raise credentials_exception
    return username

@app.post("/token")
async def login(form_data: OAuth2PasswordRequestForm = Depends()):
    # Authenticate user (implement your own logic)
    user = authenticate_user(form_data.username, form_data.password)
    if not user:
        raise HTTPException(status_code=400, detail="Incorrect username or password")
    access_token = create_access_token(data={"sub": user.username})
    return {"access_token": access_token, "token_type": "bearer"}

@app.get("/protected")
async def protected_route(current_user: str = Depends(get_current_user)):
    return {"message": f"Hello, {current_user}"}

```


2. CORS (Cross-Origin Resource Sharing) Setup

CORS is a security mechanism that allows or restricts resources on a web page to be requested from another domain outside the domain from which the resource originated.

Configure CORS in FastAPI:

```
from fastapi import FastAPI
from fastapi.middleware.cors import CORSMiddleware

app = FastAPI()

origins = [
    "http://localhost",
    "http://localhost:8080",
    "https://yourdomain.com",
]

app.add_middleware(
    CORSMiddleware,
    allow_origins=origins,
    allow_credentials=True,
    allow_methods=["*"],
    allow_headers=["*"],
)
```

3. Rate Limiting

Implementing rate limiting helps protect your API from abuse and ensures fair usage. Here's an example using the `slowapi` library:

```
from fastapi import FastAPI
from slowapi import Limiter, _rate_limit_exceeded_handler
from slowapi.util import get_remote_address
from slowapi.errors import RateLimitExceeded

limiter = Limiter(key_func=get_remote_address)
app = FastAPI()
app.state.limiter = limiter
app.add_exception_handler(RateLimitExceeded, _rate_limit_exceeded_handler)

@app.get("/")
@limiter.limit("5/minute")
```

```
async def root():  
    return {"message": "Hello World"}
```

Additional Security Measures

1. Use secure headers:

- Implement HTTP Strict Transport Security (HSTS)
- Set appropriate Content Security Policy (CSP)
- Enable X-Frame-Options to prevent clickjacking

2. Keep dependencies updated:

- Regularly update your FastAPI and other dependencies
- Use tools like `safety` to check for known vulnerabilities

3. Implement input validation:


- Leverage FastAPI's built-in request validation
- Use Pydantic models to define and validate request bodies

4. Secure your environment:

- Use a firewall to restrict access to your server
- Implement intrusion detection and prevention systems

5. Implement proper error handling:

- Avoid exposing sensitive information in error messages
- Log errors securely for debugging purposes



This diagram will help visualize the multiple layers of security that protect a FastAPI application in production.

5. Performance Optimization

Optimizing your FastAPI application for performance is crucial for handling high traffic and providing a smooth user experience. In this section, we'll explore various techniques to enhance the performance of your FastAPI application in production.

Async and Await Usage

FastAPI is built on top of Starlette and leverages Python's async capabilities. Proper use of async and await can significantly improve your application's performance, especially for I/O-bound operations.

Best Practices for Async

1. Use async functions for I/O-bound operations:

```

from fastapi import FastAPI
import httpx

```

```
app = FastAPI()

@app.get("/external-data")
async def get_external_data():
    async with httpx.AsyncClient() as client:
        response = await client.get("https://api.example.com/data")
    return response.json()
```

2. Avoid blocking operations in async functions:

```
import time
from fastapi import FastAPI
from fastapi.concurrency import run_in_threadpool

app = FastAPI()

def cpu_bound_task():
    time.sleep(1) # Simulating a CPU-bound task
    return "Task completed"

@app.get("/cpu-task")
async def handle_cpu_task():
    result = await run_in_threadpool(cpu_bound_task)
    return {"result": result}
```

Database Connection Pooling

Implementing connection pooling can significantly reduce the overhead of creating new database connections for each request.

Example using SQLAlchemy with asyncpg:

```
from fastapi import FastAPI, Depends
from sqlalchemy.ext.asyncio import create_async_engine, AsyncSession
from sqlalchemy.orm import sessionmaker

DATABASE_URL = "postgresql+asyncpg://user:password@localhost/dbname"

engine = create_async_engine(DATABASE_URL, echo=True, pool_size=20, max_overflow=10)
AsyncSessionLocal = sessionmaker(engine, class_=AsyncSession, expire_on_commit=False)
```

```
app = FastAPI()

async def get_db():
    async with AsyncSessionLocal() as session:
        yield session

@app.get("/users/{user_id}")
async def read_user(user_id: int, db: AsyncSession = Depends(get_db)):
    result = await db.execute(f"SELECT * FROM users WHERE id = {user_id}")
    user = result.fetchone()
    return {"user": user}
```

Caching Strategies

Implementing caching can dramatically improve response times for frequently accessed data.

1. In-Memory Caching with FastAPI

```
from fastapi import FastAPI
from fastapi_cache import FastAPICache
from fastapi_cache.backends.inmemory import InMemoryBackend
from fastapi_cache.decorator import cache

app = FastAPI()

@app.on_event("startup")
async def startup():
    FastAPICache.init(InMemoryBackend())

@app.get("/cached-data")
@cache(expire=60)
async def get_cached_data():
    # Simulating a slow operation
    await asyncio.sleep(2)
    return {"data": "This response is cached for 60 seconds"}
```

2. Redis Caching

For distributed systems, Redis is an excellent choice for caching:

```

from fastapi import FastAPI
from fastapi_cache import FastAPICache
from fastapi_cache.backends.redis import RedisBackend
from fastapi_cache.decorator import cache
import aioredis

app = FastAPI()

@app.on_event("startup")
async def startup():
    redis = aioredis.from_url("redis://localhost", encoding="utf8", decode_responses=True)
    FastAPICache.init(RedisBackend(redis), prefix="fastapi-cache")

@app.get("/cached-data")
@cache(expire=60)
async def get_cached_data():
    # Your data fetching logic here
    return {"data": "This response is cached in Redis for 60 seconds"}

```

Request Validation and Response Serialization Optimization

FastAPI uses Pydantic for request validation and response serialization. While this provides great benefits, it can be optimized for better performance.

1. Use `Config` class in Pydantic models to optimize:

```

from pydantic import BaseModel

class Item(BaseModel):
    name: str
    description: str
    price: float

    class Config:
        orm_mode = True
        validate_assignment = True
        allow_population_by_field_name = True

```

2. Use `response_model` parameter in route decorators to pre-compute response schemas:

```
@app.get("/items/{item_id}", response_model=Item)
async def read_item(item_id: int):
    return {"name": "Foo", "description": "An item", "price": 45.2}
```

Additional Performance Tips

1. Use UvLoop: UvLoop is a fast, drop-in replacement for the asyncio event loop.

```
import uvloop
import asyncio

asyncio.set_event_loop_policy(uvloop.EventLoopPolicy())
```

2. Implement database query optimization:

- Use database indexes effectively
- Optimize complex queries
- Use database-specific features (e.g., PostgreSQL's JSONB for complex data)

3. Implement proper logging:

- Use asynchronous logging to avoid blocking operations
- Log at appropriate levels to avoid unnecessary I/O

4. Use background tasks for time-consuming operations:

```
from fastapi import FastAPI, BackgroundTasks


app = FastAPI()

def process_item(item: dict):
    # Time-consuming operation here
    pass

@app.post("/items")
async def create_item(item: dict, background_tasks: BackgroundTasks):
```

```
background_tasks.add_task(process_item, item)
return {"message": "Item received, processing in background"}
```

5. Implement proper error handling to prevent performance degradation due to unhandled exceptions.



This diagram will help visualize how different performance optimization techniques work together in a FastAPI application to handle requests efficiently.

6. Logging and Monitoring

Proper logging and monitoring are crucial for maintaining and troubleshooting FastAPI applications in production. They help you understand your application's behavior, identify issues quickly, and make data-driven decisions for improvements.

Setting Up Proper Logging in FastAPI

FastAPI uses Python's built-in logging module. Here's how to set up logging effectively:

1. Configure logging in your main application file:


```

import logging
from fastapi import FastAPI, Request
from fastapi.logger import logger as fastapi_logger

# Configure logging
logging.basicConfig(
    level=logging.INFO,
    format='%(asctime)s - %(name)s - %(levelname)s - %(message)s',
    handlers=[
        logging.FileHandler("app.log"),
        logging.StreamHandler()
    ]
)

app = FastAPI()

@app.middleware("http")
async def log_requests(request: Request, call_next):
    logger = logging.getLogger("fastapi")
    logger.info(f"Incoming request: {request.method} {request.url}")
    response = await call_next(request)
    logger.info(f"Response status code: {response.status_code}")
    return response

@app.get("/")
async def root():
    fastapi_logger.info("Root endpoint called")
    return {"message": "Hello World"}

```

2. Use structured logging for better parsing:

```

import json
import logging
from fastapi import FastAPI, Request

class StructuredLogger(logging.Logger):
    def _log(self, level, msg, args, exc_info=None, extra=None, stack_info=False):
        if extra is None:
            extra = {}
        extra['app_name'] = 'MyFastAPIApp'
        super()._log(level, json.dumps(msg) if isinstance(msg, dict) else msg,
                      args, exc_info, stack_info, extra)

logging.setLoggerClass(StructuredLogger)

```

```
logger = logging.getLogger(__name__)

app = FastAPI()

@app.middleware("http")
async def log_structured_requests(request: Request, call_next):
    logger.info({
        "event": "request",
        "method": request.method,
        "url": str(request.url),
        "headers": dict(request.headers),
        "client": request.client.host
    })
    response = await call_next(request)
    logger.info({
        "event": "response",
        "status_code": response.status_code
    })
    return response
```

Log Rotation and Management

To prevent log files from growing too large and to manage them effectively:

1. Use a tool like `logrotate` on Linux systems:

```
/path/to/your/app.log {
    daily
    missingok
    rotate 14
    compress
    delaycompress
    notifempty
    create 0640 www-data www-data
}
```

2. Alternatively, use a Python logging handler that supports rotation:

```
import logging
from logging.handlers import RotatingFileHandler

handler = RotatingFileHandler('app.log', maxBytes=10000, backupCount=3)
```

```
logger = logging.getLogger(__name__)  
logger.addHandler(handler)
```

Monitoring Tools and Integration

1. Prometheus for metrics collection:

Get Raman Bazhanau's stories in your inbox

Join Medium for free to get updates from this writer.

Enter your email

Subscribe

Install the required libraries:

```
pip install prometheus-client starlette_exporter
```

Integrate Prometheus with FastAPI:

```
from fastapi import FastAPI  
from starlette_exporter import PrometheusMiddleware, handle_metrics  
  
app = FastAPI()  
app.add_middleware(PrometheusMiddleware)  
app.add_route("/metrics", handle_metrics)
```

2. Grafana for visualization:

- Set up Grafana to connect to your Prometheus data source
- Create dashboards to visualize your application metrics

3. ELK Stack (Elasticsearch, Logstash, Kibana) for log analysis:

- Use Filebeat to ship logs to Elasticsearch

- Use Kibana to visualize and analyze logs

Useful Metrics to Track

1. Request rate: Number of requests per second
2. Response time: Average, median, and 95th percentile
3. Error rate: Percentage of requests resulting in errors
4. CPU and memory usage
5. Database query performance
6. External API call latency
7. Active user sessions

Here's an example of how to track custom metrics using Prometheus:

```
from fastapi import FastAPI
from prometheus_client import Counter, Histogram
from starlette_exporter import PrometheusMiddleware, handle_metrics

app = FastAPI()
app.add_middleware(PrometheusMiddleware)
app.add_route("/metrics", handle_metrics)

# Define custom metrics
REQUEST_COUNT = Counter('request_count', 'Total request count')
REQUEST_LATENCY = Histogram('request_latency_seconds', 'Request latency in seconds')

@app.get("/")
async def root():
    REQUEST_COUNT.inc()
    with REQUEST_LATENCY.time():
        # Your logic here
    return {"message": "Hello World"}
```

Implementing Health Checks

Health checks are crucial for monitoring the status of your application:

```

from fastapi import FastAPI
from fastapi.responses import JSONResponse

app = FastAPI()


@app.get("/health")
async def health_check():
    # Perform checks (e.g., database connection, external services)
    all_systems_operational = True
    if all_systems_operational:
        return JSONResponse(content={"status": "healthy"}, status_code=200)
    else:
        return JSONResponse(content={"status": "unhealthy"}, status_code=503)

```

Alerting

Set up alerting based on your metrics and logs:

1. Use Prometheus Alertmanager for metric-based alerts
2. Configure alerts in Grafana for visualization-based alerting
3. Set up log-based alerts using tools like Elastic Stack's Watcher



This diagram will help visualize the comprehensive logging and monitoring setup for a FastAPI application in production, showing how different components work together to provide insights and maintain system health.

7. Containerization with Docker

Containerizing your FastAPI application with Docker provides consistency across different environments, simplifies deployment, and enhances scalability. This section will guide you through the process of containerizing your FastAPI app and best practices for production use.

Benefits of Using Docker with FastAPI

1. **Consistency:** Ensures the same environment across development, testing, and production.
2. **Isolation:** Keeps your application and its dependencies separate from the host system.
3. **Portability:** Easily move your application between different hosts or cloud providers.
4. **Scalability:** Simplifies the process of scaling your application horizontally.
5. **Version Control:** Allows versioning of your entire application environment.

Sample Dockerfile for a FastAPI Application

Here's a basic Dockerfile for a FastAPI application:

```
# Use an official Python runtime as a parent image
FROM python:3.9-slim

# Set the working directory in the container
WORKDIR /app

# Copy the current directory contents into the container at /app
COPY . /app

# Install any needed packages specified in requirements.txt
RUN pip install --no-cache-dir -r requirements.txt

# Make port 8000 available to the world outside this container
EXPOSE 8000

# Define environment variable
ENV NAME World
```

```
# Run app.py when the container launches
CMD ["uvicorn", "main:app", "--host", "0.0.0.0", "--port", "8000"]
```

Building and Running the Docker Container

To build and run your Docker container:

```
# Build the Docker image
docker build -t fastapi-app .

# Run the container
docker run -d --name myapp -p 8000:8000 fastapi-app
```

Docker Compose for Multi-Container Setups

For applications that require multiple services (e.g., FastAPI, database, Redis), use Docker Compose:

```
version: '3.8'

services:
  web:
    build: .
    ports:
      - "8000:8000"
    environment:
      - DATABASE_URL=postgresql://user:password@db:5432/dbname
    depends_on:
      - db

  db:
    image: postgres:13
    volumes:
      - postgres_data:/var/lib/postgresql/data
    environment:
      - POSTGRES_USER=user
      - POSTGRES_PASSWORD=password
      - POSTGRES_DB=dbname

volumes:
  postgres_data:
```

To run the multi-container setup:

```
docker-compose up -d
```

Best Practices for Docker in Production

1. Use specific version tags for base images:

Instead of `FROM python:3.9-slim`, use `FROM python:3.9.7-slim-buster`

2. Minimize the number of layers:

Combine RUN commands to reduce the number of layers in your image.

```
RUN apt-get update && \  
    apt-get install -y --no-install-recommends gcc && \  
    rm -rf /var/lib/apt/lists/*
```

3. Use multi-stage builds for smaller final images:

```
# Build stage  
FROM python:3.9-slim AS builder  
WORKDIR /app  
COPY requirements.txt .  
RUN pip install --user -r requirements.txt  
  
# Final stage  
FROM python:3.9-slim  
WORKDIR /app  
COPY --from=builder /root/.local /root/.local  
COPY . .  
ENV PATH=/root/.local/bin:$PATH  
CMD ["uvicorn", "main:app", "--host", "0.0.0.0", "--port", "8000"]
```

4. Don't run containers as root:

```
RUN adduser --disabled-password --gecos '' appuser
```



```
USER appuser
```

5. Use health checks:

```
HEALTHCHECK CMD curl --fail http://localhost:8000/health || exit 1
```

[Open in app](#) ↗[Sign up](#)[Sign in](#)**Medium** Search

```
ENV PYTHONDONTWRITEBYTECODE=1  
ENV PYTHONUNBUFFERED=1
```

7. Use .dockerignore file:

Create a `.dockerignore` file to exclude unnecessary files from your Docker context:

```
__pycache__  
*.pyc  
*.pyo  
*.pyd  
.git  
.env  
.venv
```

8. Implement proper logging:

Configure your application to log to stdout/stderr, which Docker can then capture.

9. Use Docker secrets for sensitive information:

Instead of hardcoding sensitive information in your Dockerfile or docker-compose.yml, use Docker secrets.

10. Regularly update base images and dependencies:

Keep your base images and Python dependencies up to date to ensure you have the latest security patches.

Optimizing Docker Image Size

1. Use Alpine-based images for even smaller footprints:

```
FROM python:3.9-alpine

# Install build dependencies
RUN apk add --no-cache gcc musl-dev linux-headers
```

2. Remove unnecessary files after installation:

```
RUN pip install --no-cache-dir -r requirements.txt && \
    rm -rf /root/.cache/pip
```

Handling Application Configuration in Docker

Use environment variables for configuration:


```
ENV APP_SETTINGS=production
ENV DATABASE_URL=postgresql://user:password@localhost/dbname
```

Then in your FastAPI app:

```
import os
from fastapi import FastAPI

app = FastAPI()

@app.get("/")
async def root():
    return {"message": f"Hello from {os.getenv('APP_SETTINGS')} environment"}
```



This diagram will help visualize the containerization process and how different components interact in a Dockerized FastAPI application setup.

8. Deployment Strategies

Deploying a FastAPI application to production requires careful planning and execution. This section will explore various deployment options, CI/CD pipelines, and strategies to ensure smooth and reliable deployments.

Deployment Options

1. Virtual Private Server (VPS)

Deploying to a VPS gives you full control over the server environment. Steps:

1. Set up a VPS with a provider like DigitalOcean, Linode, or AWS EC2.
2. Install necessary dependencies (Python, Docker, etc.).
3. Transfer your application code (via Git or SCP).
4. Set up a reverse proxy (Nginx or Apache) to handle incoming requests.
5. Configure SSL/TLS certificates (e.g., using Let's Encrypt).

6. Run your FastAPI application using a process manager like Supervisor or systemd.

Example Nginx configuration:

```
server {  
    listen 80;  
    server_name yourdomain.com;  
    return 301 https://$server_name$request_uri;  
}  
  
server {  
    listen 443 ssl;  
    server_name yourdomain.com;  
  
    ssl_certificate /path/to/fullchain.pem;  
    ssl_certificate_key /path/to/privkey.pem;  
  
    location / {  
        proxy_pass http://localhost:8000;  
        proxy_set_header Host $host;  
        proxy_set_header X-Real-IP $remote_addr;  
    }  
}
```

2. Platform as a Service (PaaS)

PaaS options like Heroku or Google App Engine can simplify deployment. Deploying to Heroku:

1. Create a Procfile:

```
web: uvicorn main:app --host 0.0.0.0 --port $PORT
```

2. Create a runtime.txt:

```
python-3.9.7
```

3. Deploy using Heroku CLI:

```
heroku create
git push heroku main
```

3. Serverless

Serverless deployment can be achieved using services like AWS Lambda with API Gateway. Using AWS Lambda:

1. Use a framework like Zappa or Mangum to adapt FastAPI for serverless.
2. Configure API Gateway to route requests to your Lambda function.
3. Deploy using AWS CLI or CloudFormation.

Example `zappa_settings.json`:

```
{
  "production": {
    "app_function": "main.app",
    "aws_region": "us-west-2",
    "project_name": "fastapi-app",
    "runtime": "python3.8",
    "s3_bucket": "zappa-fastapi-app"
  }
}
```

CI/CD Pipelines

Implementing a CI/CD pipeline automates testing and deployment processes.

GitHub Actions Example

```
name: FastAPI CI/CD

on:
  push:
    branches: [ main ]
  pull_request:
```

```
branches: [ main ]

jobs:
  test:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v2
      - name: Set up Python
        uses: actions/setup-python@v2
        with:
          python-version: 3.9
      - name: Install dependencies
        run: |
          python -m pip install --upgrade pip
          pip install -r requirements.txt
      - name: Run tests
        run: pytest

  deploy:
    needs: test
    runs-on: ubuntu-latest
    if: github.ref == 'refs/heads/main'
    steps:
      - uses: actions/checkout@v2
      - name: Deploy to Heroku
        uses: akhileshns/heroku-deploy@v3.12.12
        with:
          heroku_api_key: ${secrets.HEROKU_API_KEY}
          heroku_app_name: "your-app-name"
          heroku_email: "your-email@example.com"
```

Blue-Green Deployments

Blue-green deployment is a technique that reduces downtime and risk by running two identical production environments called Blue and Green. Steps:

1. Set up two identical environments (Blue and Green).
2. Route all traffic to the Blue environment.
3. Deploy new version to the Green environment.
4. Test the Green environment.
5. Switch traffic from Blue to Green.
6. Keep Blue as a quick rollback option.

Example using Nginx for traffic routing:

```
upstream backend {  
    server blue-server;  
    # server green-server backup;  
}  
  
server {  
    listen 80;  
    server_name example.com;  
  
    location / {  
        proxy_pass http://backend;  
    }  
}
```

To switch to Green, update the Nginx configuration and reload:

```
upstream backend {  
    server green-server;  
    # server blue-server backup;  
}
```

Rolling Updates

Rolling updates involve gradually replacing instances of the old version with the new version. Steps:

1. Deploy new version alongside the old version.
2. Gradually shift traffic to the new version.
3. Monitor for any issues.
4. If successful, complete the rollout; if not, rollback.

Example using Docker Swarm:

```
version: '3'
services:
  web:
    image: your-fastapi-app:latest
    deploy:
      replicas: 4
      update_config:
        parallelism: 2
        order: rolling-update
```

Update the service:

```
docker service update --image your-fastapi-app:new-version your-service-name
```

Deployment Scripts

Creating deployment scripts can automate and standardize the deployment process.

Example deployment script:

```
#!/bin/bash

set -e

echo "Deploying FastAPI application..."


# Pull latest changes
git pull origin main

# Install dependencies
pip install -r requirements.txt

# Run migrations
alembic upgrade head

# Restart the application
sudo systemctl restart fastapi-app

echo "Deployment completed successfully!"
```

This diagram will help visualize the complex process of deploying a FastAPI application to production, including various strategies and safety mechanisms.

9. Scaling FastAPI Applications

As your FastAPI application grows in popularity, you'll need to implement scaling strategies to handle increased traffic and maintain performance. This section will explore various methods to scale your FastAPI application effectively.

Horizontal vs. Vertical Scaling

Before diving into specific strategies, it's important to understand the two main types of scaling:

1. Vertical Scaling (Scaling Up): Increasing the resources (CPU, RAM) of existing servers.
2. Horizontal Scaling (Scaling Out): Adding more servers to distribute the load.

FastAPI, being built on asynchronous principles, is well-suited for horizontal scaling, which we'll focus on in this section.

Load Balancing

Load balancing is crucial for distributing incoming requests across multiple instances of your FastAPI application.

Nginx as a Load Balancer

Here's an example Nginx configuration for load balancing:

```
http {
    upstream fastapi_servers {
        server 127.0.0.1:8000;
        server 127.0.0.1:8001;
        server 127.0.0.1:8002;
    }

    server {
        listen 80;
        server_name example.com;

        location / {
            proxy_pass http://fastapi_servers;
            proxy_set_header Host $host;
            proxy_set_header X-Real-IP $remote_addr;
        }
    }
}
```

Using Docker Swarm for Load Balancing

Docker Swarm provides built-in load balancing for containerized applications:

```
version: '3'
services:
  fastapi_app:
    image: your-fastapi-app:latest
    deploy:
      replicas: 5
      update_config:
        parallelism: 2
        order: rolling-update
    ports:
      - "8000:8000"
```

Database Scaling

As your application scales, your database can become a bottleneck. Here are some strategies to scale your database:

1. **Read Replicas:** Create read-only copies of your database to offload read operations.
2. **Sharding:** Distribute data across multiple databases based on a shard key.
3. **Connection Pooling:** Use a connection pool to manage database connections efficiently.

Example of connection pooling with SQLAlchemy:

```
from sqlalchemy import create_engine
from sqlalchemy.orm import sessionmaker
from fastapi import FastAPI, Depends

DATABASE_URL = "postgresql://user:password@localhost/dbname"

engine = create_engine(DATABASE_URL, pool_size=20, max_overflow=0)
SessionLocal = sessionmaker(autocommit=False, autoflush=False, bind=engine)

app = FastAPI()

def get_db():
    db = SessionLocal()
    try:
        yield db
    finally:
        db.close()

@app.get("/users/{user_id}")
def read_user(user_id: int, db: Session = Depends(get_db)):
    user = db.query(User).filter(User.id == user_id).first()
    return user
```

Caching

Implementing caching can significantly reduce the load on your application and database.

Using Redis for Caching

```
from fastapi import FastAPI
from fastapi_cache import FastAPICache
```

```
from fastapi_cache.backends.redis import RedisBackend
from fastapi_cache.decorator import cache
import aioredis

app = FastAPI()

@app.on_event("startup")
async def startup():
    redis = aioredis.from_url("redis://localhost", encoding="utf8", decode_responses=True)
    FastAPICache.init(RedisBackend(redis), prefix="fastapi-cache")

@app.get("/expensive-operation")
@cache(expire=60)
async def expensive_operation():
    # Perform expensive operation here
    return {"result": "Expensive operation result"}
```

Asynchronous Task Processing

For long-running tasks, use asynchronous task queues to offload work from the main application.

Using Celery with FastAPI

```
from fastapi import FastAPI
from celery import Celery

app = FastAPI()
celery = Celery("tasks", broker="redis://localhost:6379/0")

@celery.task
def process_data(data):
    # Process data asynchronously
    pass

@app.post("/process")
async def process_endpoint(data: dict):
    task = process_data.delay(data)
    return {"task_id": task.id}

@app.get("/task/{task_id}")
async def get_task_status(task_id: str):
```

```
task = process_data.AsyncResult(task_id)
return {"status": task.status, "result": task.result}
```

Serverless Scaling

Serverless platforms like AWS Lambda can automatically scale your application based on incoming requests.

Example using Mangum to adapt FastAPI for AWS Lambda:

```
from fastapi import FastAPI
from mangum import Mangum

app = FastAPI()

@app.get("/")
async def root():
    return {"message": "Hello World"}

handler = Mangum(app)
```

Deploy this to AWS Lambda and configure API Gateway to route requests to your function.

Monitoring and Auto-scaling

Implement monitoring and auto-scaling to automatically adjust resources based on traffic.

Using AWS Auto Scaling

1. Create an Amazon EC2 Auto Scaling group for your FastAPI instances.
2. Set up CloudWatch alarms to monitor metrics like CPU utilization or request count.
3. Configure Auto Scaling policies to add or remove instances based on these alarms.

Example Auto Scaling policy (AWS CLI):

```
aws autoscaling put-scaling-policy \  
  --auto-scaling-group-name my-fastapi-asg \  
  --policy-name cpu-policy \  
  --policy-type TargetTrackingScaling \  
  --target-tracking-configuration file://config.json
```

Where config.json contains:


```
{  
  "TargetValue": 70.0,  
  "PredefinedMetricSpecification": {  
    "PredefinedMetricType": "ASGAverageCPUUtilization"  
  }  
}
```

Content Delivery Network (CDN)

Utilize a CDN to cache and serve static content from locations closer to users.


Example using Cloudflare with FastAPI:

1. Sign up for a Cloudflare account and add your domain.
2. Update your domain's nameservers to Cloudflare's.
3. Enable caching for static content in Cloudflare's dashboard.
4. Optionally, use Cloudflare Workers to add edge computing capabilities.



1. **Client Requests:** The entry point for user interactions.
2. **CDN (Content Delivery Network):** Routes requests and caches static content.
3. **Load Balancer:** Distributes incoming requests across multiple FastAPI instances.
4. **Multiple FastAPI Application Instances:** Handle incoming requests.
5. **Auto Scaling Group:** Adjusts the number of FastAPI instances based on demand.
6. **Monitoring System:** Observes system metrics and triggers scaling actions.

The flow shows how requests move from clients through the CDN and load balancer to the FastAPI instances, and how the auto-scaling group and monitoring system work together to adjust resources.



1. **FastAPI Instances:** Represent the application servers.
2. **Caching Layer (Redis):** Stores frequently accessed data.
3. **Main Database:** Handles write operations.
4. **Read Replicas:** Distribute read operations.
5. **Task Queue (Celery):** Manages asynchronous tasks.
6. **Worker Nodes:** Process tasks from the queue.
7. **Serverless Functions:** Handle specific tasks or spikes in traffic.

This diagram shows how FastAPI instances interact with various data storage systems (cache, main database, read replicas) and how asynchronous tasks are processed using a queue and worker nodes. It also illustrates the option of using serverless functions for specific tasks.

10. Conclusion and Best Practices

As we conclude this comprehensive guide on preparing FastAPI for production, let's summarize the key points and provide a set of best practices to ensure your FastAPI application is production-ready, performant, and scalable.

Key Takeaways

1. **Configuration Management:** Properly configure your FastAPI application for different environments using environment variables and configuration management tools.
2. **ASGI Servers:** Choose the right ASGI server (Uvicorn, Hypercorn, or Gunicorn with Uvicorn workers) based on your specific needs.
3. **Security:** Implement robust security measures, including HTTPS, authentication, authorization, and proper error handling.
4. **Performance Optimization:** Utilize FastAPI's async capabilities, implement caching, and optimize database queries for better performance.
5. **Logging and Monitoring:** Set up comprehensive logging and monitoring to gain insights into your application's behavior and quickly identify issues.
6. **Containerization:** Use Docker to containerize your FastAPI application for consistency across environments and easier deployment.
7. **Deployment Strategies:** Implement CI/CD pipelines and use deployment strategies like blue-green deployments or rolling updates to minimize downtime.
8. **Scaling:** Prepare your application for scaling by implementing load balancing, database scaling strategies, and considering serverless options.

Best Practices Checklist

To ensure your FastAPI application is production-ready, follow this checklist:

1. Security

- Enable HTTPS using SSL/TLS certificates
- Implement proper authentication and authorization
- Set up CORS correctly
- Use secure headers (HSTS, CSP, X-Frame-Options)
- Implement rate limiting to prevent abuse
- Regularly update dependencies to patch security vulnerabilities

2. Performance

- Use asynchronous programming where appropriate
- Implement caching for frequently accessed data
- Optimize database queries and use connection pooling
- Utilize background tasks for time-consuming operations
- Configure your ASGI server for optimal performance

3. Reliability

- Implement proper error handling and logging
- Set up health check endpoints
- Use database migrations for schema changes
- Implement retry mechanisms for external service calls
- Set up automated backups for your database

4. Scalability

- Design your application to be stateless
- Implement horizontal scaling with load balancing
- Use caching services like Redis to reduce database load
- Consider using a CDN for static assets
- Implement database scaling strategies (read replicas, sharding)

5. Monitoring and Logging

- Set up centralized logging
- Implement application performance monitoring (APM)
- Set up alerts for critical errors and performance issues
- Use structured logging for easier parsing and analysis

- Monitor key metrics like response times, error rates, and resource usage

6. Deployment and Operations

- Use containerization (Docker) for consistent environments
- Implement a CI/CD pipeline for automated testing and deployment
- Use infrastructure as code (IaC) for managing your infrastructure
- Implement blue-green or rolling update deployment strategies
- Set up automated scaling based on load

7. Documentation and Maintenance

- Keep API documentation up-to-date (use FastAPI's automatic docs)
- Document deployment procedures and runbooks
- Maintain a changelog for tracking changes
- Regularly review and update your application's dependencies
- Conduct periodic security audits and penetration testing

Final Thoughts

Preparing a FastAPI application for production is a multifaceted process that requires attention to various aspects of software development and operations. By following the practices and strategies outlined in this guide, you can create a robust, scalable, and maintainable FastAPI application that performs well under real-world conditions.

Remember that production readiness is an ongoing process. Continuously monitor your application, stay updated with the latest best practices and security recommendations, and be prepared to adapt your infrastructure and codebase as your application grows and evolves.

Lastly, leverage the FastAPI community and ecosystem. Stay engaged with the community forums, contribute to open-source projects, and don't hesitate to seek help when faced with challenges. The collective knowledge and experience of the community can be an invaluable resource as you navigate the complexities of running FastAPI in production.

Let's connect!
[LinkedIn](#)

Python

Fastapi

Deployment

Coding

Software Development



Follow

Written by Raman Bazhanau

648 followers · 2 following

Software Engineer

Responses (17)



Write a response

What are your thoughts?



Paulo Coutinho
Oct 21, 2024



I have a boilerplate for FastAPI ready for prod here: <https://github.com/paulocoutinhox/fastapi-app>

It is free and open source. Feel free to make suggestions.



23



1 reply

[Reply](#)



Sorower Hossan

Feb 18, 2025



Here's an example of how to track custom metrics using Prometheus:

what to do when I use gunicorn with multiple workers?

it saves the logs to different workers, so each time i get different numbers :(



2

[Reply](#)



Haley she/her

Jul 24, 2025



I feat this may be the most helpful article I've ever read in my life.




1

[Reply](#)

See all responses

More from Raman Bazhanau



Mastering SQLAlchemy: A Comprehensive Guide for Python Developers

Enter SQLAlchemy, one of the most powerful and flexible ORMs available for Python.

Nov 6, 2024  267  5



Advanced Type Annotations in Python: Part 2

In the first part of our deep dive into Python's type annotations, we explored the foundational tools like Generic, TypeVar, and TypedDict...

Sep 19, 2023  264  1





Mastering Advanced OOP Concepts in Python: Advanced Class Features

Welcome to this deep dive into the world of advanced Object-Oriented Programming (OOP) in Python. Whether you're an intermediate Python...

Dec 14, 2023  475  2



Python Memory Management: Best Practices for Performance

A Comprehensive Guide to Mastering Memory Efficiency and Boosting Performance in Python Applications

Oct 26, 2024  233  3



See all from Raman Bazhanau

Recommended from Medium

 Hash Block

FastAPI at 200K Requests Per Minute: Horizontal Scaling with Gunicorn and Kubernetes HPA

How I Built a Resilient FastAPI Architecture to Handle Traffic Spikes with Kubernetes Autoscaling and Gunicorn Workers

 Jul 29, 2025  57



 In Python in Plain English by Sandun Lakshan

8 FastAPI Best Practices Every Python Developer Should Know

FastAPI has become one of the most popular frameworks for building APIs in Python. Its speed, modern design, and strong developer...

✦ Sep 10, 2025 🖱 31



 In Level Up Coding by Arslan Qutab

7 FastAPI Extensions So Powerful, I Stopped Using Flask Forever

From authentication to background tasks here's what made FastAPI my go-to for every new project.



Sep 7, 2025



659



8



In Artificial Intelligence in Plain English by Arijit Deb

RAG Pipeline using FastAPI-Celery-Redis

RAG pipelines are everywhere now — powering chatbots, assistants, and enterprise search. Instead of revisiting what RAG is, let's talk...



Nov 4, 2025



40



1



Codastra

FastAPI + Rate Limiting with Redis: Fair-Use APIs Without User Rage

How to protect your backend, keep paying customers happy, and avoid “your API sucks” tweets.

★ Dec 2, 2025 🖱 45



○ Jahel

Production ready event driven architecture with fastapi + docker + AWS: PART 1

This guide walks you through building a production-ready FastAPI application with event-driven architecture, starting with a single...

Dec 28, 2025



See more recommendations